

Classification Models

November 7, 2021

```
[1]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_validate
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.model_selection import StratifiedKFold
from sklearn import ensemble
import xgboost as xgb
from sklearn.metrics import mean_squared_error
```

Import Dataset

```
[2]: df =pd.read_csv ('/home/jovyan/Onlineshoppersdata(1).csv')
```

1 Classification Algorithms with imbalanced dataset

1.1 Random Forest

1.1.1 Separate the data into Train and Test

```
[3]: x = df.drop(['Revenue'],axis=1).values    # independant features
      y = df['Revenue'].values                # dependant_
      ↪variable

      x_trainrf, x_testrf, y_trainrf, y_testrf = train_test_split(x, y, test_size=0.
      ↪2, random_state=42)
```

```
[4]: forest = RandomForestClassifier()

      forest.fit(x_trainrf,y_trainrf)

      y_predrf =forest.predict(x_testrf)
```

```
[5]: accuracy_forest = (metrics.accuracy_score(y_testrf,y_predrf)*100)
      print('Acurracy is ', str(round(accuracy_forest,2)))
```

Acurracy is 89.05

```
[6]: rmse = np.sqrt(mean_squared_error(y_testrf, y_predrf))
      print("RMSE (root-mean-square error): %f" % (rmse))
```

RMSE (root-mean-square error): 0.330891

```
[7]: # View confusion matrix for test data and predictions
      matrix_rf = confusion_matrix(y_testrf, y_predrf)
      print(matrix_rf )
```

```
[[1974  81]
 [ 189 222]]
```


```
[8]: # Get and reshape confusion matrix data

      matrix_rforest = matrix_rf.astype('float') / matrix_rf.sum(axis=1)[: , np.
      ↪newaxis]

      # Build the plot
      plt.figure(figsize=(10,5))
      sns.set(font_scale=1.4)
      sns.heatmap(matrix_rforest, annot=True, annot_kws={'size':10},
                  cmap=plt.cm.Greens, linewidths=0.2)

      # Add labels to the plot
      class_names = ['Yes', 'No']
      tick_marks = np.arange(len(class_names))
      tick_marks2 = tick_marks + 0.5
      plt.xticks(tick_marks, class_names, rotation=0,)
      plt.yticks(tick_marks2, class_names, rotation=0, )
```

```
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion Matrix for Random Forest Model')
plt.savefig('output14.png', dpi=300, bbox_inches='tight')
plt.show()
```



```
[9]: # View the classification report for test data and predictions
print(classification_report(y_testrf, y_predrf))
```

	precision	recall	f1-score	support
0	0.91	0.96	0.94	2055
1	0.73	0.54	0.62	411
accuracy			0.89	2466
macro avg	0.82	0.75	0.78	2466
weighted avg	0.88	0.89	0.88	2466

1.2 K-Fold Cross Validation for Random Forest

Using Cross Validate

```
[10]: x = df.drop(['Revenue'],axis=1).values # independant features
```

```

y = df['Revenue'].values # dependant_
↪variable

model = RandomForestClassifier(random_state=42)
cv = cross_validate(model, x, y, cv=10)
test_score = (cv['test_score'])*100
print('The test scores are:, ' + str(test_score))
print('The average test score is :' + str(test_score.mean()))

```

The test scores are:, [88.24006488 90.51094891 93.18734793 88.40227088
89.6188159 87.10462287
87.34793187 88.80778589 89.21330089 86.13138686]
The average test score is :88.85644768856449

```

[11]: scores = cross_val_score(model, x, y, cv=10)
scores
print("%.2f accuracy with a standard deviation of %.2f" % (scores.mean(),_
↪scores.std()))

```

0.89 accuracy with a standard deviation of 0.02

Using Kfold to separate the data

```

[12]: X = df.drop(['Revenue'],axis=1).values # independant features
y = df['Revenue'].values # dependant_
↪variable

kf = KFold(n_splits=10,random_state=42,shuffle=True)
kf.get_n_splits(X)
i=1
print("confusion matrix:")
for train_index, test_index in kf.split(X):

    X_train, X_test = X[train_index], X[test_index]

    y_train, y_test = y[train_index], y[test_index]

    model.fit(X_train, y_train)

    print (i,"\n",confusion_matrix(y_test, model.predict(X_test)))
    i=i+1
    print(10* '#')

```

confusion matrix:

```

1
[[987  43]
 [ 95 108]]
#####

```

```

2
[[989 36]
 [ 91 117]]
#####
3
[[1029 40]
 [ 67 97]]
#####
4
[[1007 39]
 [ 79 108]]
#####
5
[[1013 39]
 [ 76 105]]
#####
6
[[1011 28]
 [ 75 119]]
#####
7
[[1010 37]
 [ 79 107]]
#####
8
[[984 46]
 [ 89 114]]
#####
9
[[1007 28]
 [ 72 126]]
#####
10
[[1012 37]
 [ 92 92]]
#####

```

Using StratifiedKFold

```

[13]: X = df.drop(['Revenue'],axis=1).values    # independant features
      y = df['Revenue'].values                  # dependant
      ↪variable

kf = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

cnt = 1
# split() method generate indices to split data into training and test set.
for train_index, test_index in kf.split(X, y):

```

```

    print(f'Fold:{cnt}, Train set: {len(train_index)}, Test set:
    ↳{len(test_index)}')
    cnt+=1

score = cross_val_score(ensemble.RandomForestClassifier(random_state= 42), X,
    ↳y, cv= kf, scoring="accuracy")
print(f'Scores for each fold are: {score}')
print(f'Average score: {"{: .2f}".format(score.mean())}')

```

```

Fold:1, Train set: 11097, Test set:1233
Fold:2, Train set: 11097, Test set:1233
Fold:3, Train set: 11097, Test set:1233
Fold:4, Train set: 11097, Test set:1233
Fold:5, Train set: 11097, Test set:1233
Fold:6, Train set: 11097, Test set:1233
Fold:7, Train set: 11097, Test set:1233
Fold:8, Train set: 11097, Test set:1233
Fold:9, Train set: 11097, Test set:1233
Fold:10, Train set: 11097, Test set:1233
Scores for each fold are: [0.90186537 0.9026764  0.90024331 0.89699919
0.90510949 0.90348743
 0.9026764  0.9107867  0.89213301 0.9026764 ]
Average score: 0.90

```

After applying K-fold to Random Forest we can observe that the average score remains at 90%

1.3 XGBoost

```

[14]: x = df.drop(['Revenue'],axis=1).values    # independant features
      y = df['Revenue'].values                  # dependant
      ↳variable

data_dmatrix = xgb.DMatrix(data=x,label=y)
X_trainxg, X_testxg, y_trainxg, y_testxg = train_test_split(x, y, test_size=0.
    ↳2, random_state=42)

xg_model = xgb.XGBClassifier()
print(xg_model)

```

```

XGBClassifier(base_score=None, booster=None, colsample_bylevel=None,
               colsample_bynode=None, colsample_bytree=None,
               enable_categorical=False, gamma=None, gpu_id=None,
               importance_type=None, interaction_constraints=None,
               learning_rate=None, max_delta_step=None, max_depth=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               n_estimators=100, n_jobs=None, num_parallel_tree=None,

```

```

        predictor=None, random_state=None, reg_alpha=None,
        reg_lambda=None, scale_pos_weight=None, subsample=None,
        tree_method=None, validate_parameters=None, verbosity=None)

```

```

[15]: xg_model.fit(X_trainxg,y_trainxg)

preds_xg = xg_model.predict(X_testxg)# make predictions for test data

predictions = [round(value) for value in preds_xg]

```

[22:30:20] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
warnings.warn(label_encoder_deprecation_msg, UserWarning)

```

[16]: accuracy_xg = (metrics.accuracy_score(y_testxg,preds_xg)*100)
print('Accuracy is ', str(round(accuracy_forest,2)))

```

Accuracy is 89.05

```

[17]: rmse = np.sqrt(mean_squared_error(y_testxg, preds_xg))
print("RMSE (root-mean-square error): %f" % (rmse))

```

RMSE (root-mean-square error): 0.338165

```

[24]: # View confusion matrix for test data and predictions
matrix_xgboost=confusion_matrix(y_testxg, preds_xg)
print(matrix_xgboost)

```

```

[[1958  97]
 [ 185 226]]

```

```

[25]: # Get and reshape confusion matrix data
matrix_xg = matrix_xgboost.astype('float') / matrix_xgboost.sum(axis=1)[: , np.
→newaxis]


# Build the plot
plt.figure(figsize=(10,5))
sns.set(font_scale=1.4)
sns.heatmap(matrix_xg, annot=True, annot_kws={'size':10},
            cmap=plt.cm.Greens, linewidths=0.2)

```

```

# Add labels to the plot
class_names = ['Yes', 'No']
tick_marks = np.arange(len(class_names))
tick_marks2 = tick_marks + 0.5
plt.xticks(tick_marks, class_names, rotation=0,)
plt.yticks(tick_marks2, class_names, rotation=0, )
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion Matrix for XGBoost')
plt.savefig('output15.png', dpi=300, bbox_inches='tight')
plt.show()

```



```
[26]: print(classification_report(y_testxg, preds_xg))
```

	precision	recall	f1-score	support
0	0.91	0.95	0.93	2055
1	0.70	0.55	0.62	411
accuracy			0.89	2466
macro avg	0.81	0.75	0.77	2466
weighted avg	0.88	0.89	0.88	2466

1.4 XGBoost Model With k-Fold Cross Validation

Using Kfold

```
[33]: x = df.drop(['Revenue'],axis=1).values    # independant features
      y = df['Revenue'].values                  # dependant
      ↪variable

      # CV model
      model = xgb.XGBClassifier(use_label_encoder=False,eval_metric='mlogloss' )
      kfold = KFold(n_splits=10, random_state=42, shuffle=True)
      results = cross_val_score(model, X, y, cv=kfold)
      print("Accuracy: %.2f%% (%.2f%%)" % (results.mean()*100, results.std()*100))
```

Accuracy: 89.70% (0.87%)

Using stratified cross validation

```
[36]: X = df.drop(['Revenue'],axis=1).values    # independant features
      Y = df['Revenue'].values                  # dependant
      ↪variable

      # CV model
      model = xgb.XGBClassifier(use_label_encoder=False,eval_metric='mlogloss' )
      kfold = StratifiedKFold(n_splits=10, random_state=7,shuffle=True)
      results = cross_val_score(model, X, Y, cv=kfold)
      print("The Accuracy is: %.2f%% (%.2f%%)" % (results.mean()*100, results.
      ↪std()*100))
```

The Accuracy is: 89.75% (0.56%)

After applying K-fold to the XGBoost model we can observe that the average score remains at 88-89%

```
[ ]:
```