

Classification Models

November 19, 2021

```
[1]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_validate
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.model_selection import StratifiedKFold
from sklearn import ensemble
import xgboost as xgb
from sklearn.metrics import mean_squared_error
```

Import Dataset

```
[2]: df =pd.read_csv ('/home/jovyan/Onlineshoppersdata(1).csv')
```

1 Classification Algorithms with imbalanced dataset

1.1 Random Forest

1.1.1 Separate the data into Train and Test

```
[3]: x = df.drop(['Revenue'],axis=1).values    # independant features
      y = df['Revenue'].values                # dependant_
      ↪variable

      x_trainrf, x_testrf, y_trainrf, y_testrf = train_test_split(x, y, test_size=0.
      ↪2, random_state=42)
```

```
[4]: forest = RandomForestClassifier()

      forest.fit(x_trainrf,y_trainrf)

      y_predrf =forest.predict(x_testrf)
```

```
[5]: accuracy_forest = (metrics.accuracy_score(y_testrf,y_predrf)*100)
      print('Acurracy is ', str(round(accuracy_forest,2)))
```

Acurracy is 89.13

```
[6]: rmse = np.sqrt(mean_squared_error(y_testrf, y_predrf))
      print("RMSE (root-mean-square error): %f" % (rmse))
```

RMSE (root-mean-square error): 0.329663

```
[7]: # View confusion matrix for test data and predictions
      matrix_rf = confusion_matrix(y_testrf, y_predrf)
      print(matrix_rf )
```

```
[[1982  73]
 [ 195 216]]
```


```
[8]: # Get and reshape confusion matrix data

      matrix_rforest = matrix_rf.astype('float') / matrix_rf.sum(axis=1)[: , np.
      ↪newaxis]

      # Build the plot
      plt.figure(figsize=(10,5))
      sns.set(font_scale=1.4)
      sns.heatmap(matrix_rforest, annot=True, annot_kws={'size':10},
                  cmap=plt.cm.Greens, linewidths=0.2)

      # Add labels to the plot
      class_names = ['Yes', 'No']
      tick_marks = np.arange(len(class_names))
      tick_marks2 = tick_marks + 0.5
      plt.xticks(tick_marks, class_names, rotation=0,)
      plt.yticks(tick_marks2, class_names, rotation=0, )
```

```
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion Matrix for Random Forest Model')
plt.savefig('output14.png', dpi=300, bbox_inches='tight')
plt.show()
```



```
[9]: # View the classification report for test data and predictions
print(classification_report(y_testrf, y_predrf))
```

	precision	recall	f1-score	support
0	0.91	0.96	0.94	2055
1	0.75	0.53	0.62	411
accuracy			0.89	2466
macro avg	0.83	0.75	0.78	2466
weighted avg	0.88	0.89	0.88	2466

1.2 K-Fold Cross Validation for Random Forest

```
[34]: X = df.drop(['Revenue'],axis=1).values # independant features
      y = df['Revenue'].values # dependant variable
      ↪variable
```

```

model= RandomForestClassifier()
kf = KFold(n_splits=10,random_state=42,shuffle=True)
kf.get_n_splits(X)
i=1
print("confusion matrix:")
for train_index, test_index in kf.split(X):

    X_train, X_test = X[train_index], X[test_index]

    y_train, y_test_kfold = y[train_index], y[test_index]

    model.fit(X_train, y_train)

    print (i,"\n",confusion_matrix(y_test, model.predict(X_test)))
    i=i+1
    print(10* '#')

```

confusion matrix:

```

1
[[925 124]
 [154  30]]
#####
2
[[930 119]
 [151  33]]
#####
3
[[928 121]
 [160  24]]
#####
4
[[931 118]
 [162  22]]
#####
5
[[932 117]
 [161  23]]
#####
6
[[928 121]
 [165  19]]
#####
7
[[933 116]
 [153  31]]
#####

```

```

8
[[908 141]
 [159 25]]
#####
9
[[929 120]
 [151 33]]
#####
10
[[1011 38]
 [ 91 93]]
#####

```

```
[35]: y_pred_kfold=model.predict(X_test)
```

```
[36]: # View the classification report for test data and predictions
print(classification_report(y_test_kfold, y_pred_kfold))
```

	precision	recall	f1-score	support
0	0.92	0.96	0.94	1049
1	0.71	0.51	0.59	184
accuracy			0.90	1233
macro avg	0.81	0.73	0.77	1233
weighted avg	0.89	0.90	0.89	1233

```
[37]: scores = cross_val_score(model, X, y, cv=kf)
scores
print("%.2f accuracy with a standard deviation of %.2f" % (scores.mean(),
↪scores.std()))
```

0.90 accuracy with a standard deviation of 0.01

After applying K-fold to Random Forest we can observe that the average score remains at 90%

1.3 XGBoost

```
[14]: x = df.drop(['Revenue'],axis=1).values # independant features
y = df['Revenue'].values # dependant
↪variable

data_dmatrix = xgb.DMatrix(data=x,label=y)
X_trainxg, X_testxg, y_trainxg, y_testxg = train_test_split(x, y, test_size=0.
↪2, random_state=42)
```

```
xg_model = xgb.XGBClassifier()
print(xg_model)
```

```
XGBClassifier(base_score=None, booster=None, colsample_bylevel=None,
              colsample_bynode=None, colsample_bytree=None,
              enable_categorical=False, gamma=None, gpu_id=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=None, max_delta_step=None, max_depth=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              n_estimators=100, n_jobs=None, num_parallel_tree=None,
              objective='binary:logistic', predictor=None, random_state=None,
              reg_alpha=None, reg_lambda=None, scale_pos_weight=None,
              subsample=None, tree_method=None, use_label_encoder=True,
              validate_parameters=None, verbosity=None)
```

```
[15]: xg_model.fit(X_trainxg,y_trainxg)

preds_xg = xg_model.predict(X_testxg)# make predictions for test data

predictions = [round(value) for value in preds_xg]
```

[17:27:13] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
[16]: accuracy_xg = (metrics.accuracy_score(y_testxg,preds_xg)*100)
print('Accuracy is ', str(round(accuracy_forest,2)))
```

Accuracy is 89.13

```
[17]: rmse = np.sqrt(mean_squared_error(y_testxg, preds_xg))
print("RMSE (root-mean-square error): %f" % (rmse))
```

RMSE (root-mean-square error): 0.338165


```
[18]: # View confusion matrix for test data and predictions
matrix_xgboost=confusion_matrix(y_testxg, preds_xg)
print(matrix_xgboost)
```

```
[[1958   97]
 [ 185  226]]
```

```
[19]: # Get and reshape confusion matrix data
matrix_xg = matrix_xgboost.astype('float') / matrix_xgboost.sum(axis=1)[: , np.
      ↪newaxis]

# Build the plot
plt.figure(figsize=(10,5))
sns.set(font_scale=1.4)
sns.heatmap(matrix_xg, annot=True, annot_kws={'size':10},
            cmap=plt.cm.Greens, linewidths=0.2)

# Add labels to the plot
class_names = ['Yes', 'No']
tick_marks = np.arange(len(class_names))
tick_marks2 = tick_marks + 0.5
plt.xticks(tick_marks, class_names, rotation=0,)
plt.yticks(tick_marks2, class_names, rotation=0, )
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.title('Confusion Matrix for XGBoost')
plt.savefig('output15.png', dpi=300, bbox_inches='tight')
plt.show()
```



```
[20]: print(classification_report(y_testxg, preds_xg))
```

	precision	recall	f1-score	support
0	0.91	0.95	0.93	2055
1	0.70	0.55	0.62	411
accuracy			0.89	2466
macro avg	0.81	0.75	0.77	2466
weighted avg	0.88	0.89	0.88	2466

1.4 XGBoost Model With k-Fold Cross Validation

Using Kfold

```
[40]: x = df.drop(['Revenue'],axis=1).values    # independant features
      y = df['Revenue'].values                # dependant variable
      ↪variable

# CV model
modelxg = xgb.XGBClassifier()
kf = KFold(n_splits=10,random_state=42,shuffle=True)
kf.get_n_splits(X)
i=1
print("confusion matrix:")
for train_index, test_index in kf.split(X):

    X_train, X_test = X[train_index], X[test_index]

    y_train, y_test_xg = y[train_index], y[test_index]

    modelxg.fit(X_train, y_train)

    print (i,"\n",confusion_matrix(y_test, model.predict(X_test)))
    i=i+1
print(10* '#')
```

confusion matrix:

[17:45:02] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode


```

your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
    warnings.warn(label_encoder_deprecation_msg, UserWarning)

1
    [[880 169]
     [150 34]]
#####
[17:45:04] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The
use of label encoder in XGBClassifier is deprecated and will be removed in a
future release. To remove this warning, do the following: 1) Pass option
use_label_encoder=False when constructing XGBClassifier object; and 2) Encode
your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
    warnings.warn(label_encoder_deprecation_msg, UserWarning)

2
    [[881 168]
     [144 40]]
#####
[17:45:06] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The
use of label encoder in XGBClassifier is deprecated and will be removed in a
future release. To remove this warning, do the following: 1) Pass option
use_label_encoder=False when constructing XGBClassifier object; and 2) Encode
your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
    warnings.warn(label_encoder_deprecation_msg, UserWarning)

3
    [[910 139]
     [159 25]]
#####
[17:45:08] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The
use of label encoder in XGBClassifier is deprecated and will be removed in a
future release. To remove this warning, do the following: 1) Pass option
use_label_encoder=False when constructing XGBClassifier object; and 2) Encode
your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
    warnings.warn(label_encoder_deprecation_msg, UserWarning)

```

4

```
[[890 159]
[156 28]]
#####
[17:45:10] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The
use of label encoder in XGBClassifier is deprecated and will be removed in a
future release. To remove this warning, do the following: 1) Pass option
use_label_encoder=False when constructing XGBClassifier object; and 2) Encode
your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
  warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

5

```
[[895 154]
[157 27]]
#####
[17:45:12] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The
use of label encoder in XGBClassifier is deprecated and will be removed in a
future release. To remove this warning, do the following: 1) Pass option
use_label_encoder=False when constructing XGBClassifier object; and 2) Encode
your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
  warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

6

```
[[879 170]
[160 24]]
#####
[17:45:14] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The
use of label encoder in XGBClassifier is deprecated and will be removed in a
future release. To remove this warning, do the following: 1) Pass option
use_label_encoder=False when constructing XGBClassifier object; and 2) Encode
your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
  warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

7

```
[[895 154]
[152 32]]
```

```

#####
[17:45:16] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The
use of label encoder in XGBClassifier is deprecated and will be removed in a
future release. To remove this warning, do the following: 1) Pass option
use_label_encoder=False when constructing XGBClassifier object; and 2) Encode
your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
  warnings.warn(label_encoder_deprecation_msg, UserWarning)

8
[[879 170]
 [151  33]]
#####
[17:45:18] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The
use of label encoder in XGBClassifier is deprecated and will be removed in a
future release. To remove this warning, do the following: 1) Pass option
use_label_encoder=False when constructing XGBClassifier object; and 2) Encode
your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
  warnings.warn(label_encoder_deprecation_msg, UserWarning)

9
[[892 157]
 [144  40]]
#####
[17:45:20] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the
default evaluation metric used with the objective 'binary:logistic' was changed
from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore
the old behavior.

/opt/conda/lib/python3.7/site-packages/xgboost/sklearn.py:1224: UserWarning: The
use of label encoder in XGBClassifier is deprecated and will be removed in a
future release. To remove this warning, do the following: 1) Pass option
use_label_encoder=False when constructing XGBClassifier object; and 2) Encode
your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
  warnings.warn(label_encoder_deprecation_msg, UserWarning)

10
[[1011  38]
 [  91  93]]
#####

```

```
[45]: y_pred_xg=modelxg.predict(X_test)
      print(classification_report(y_test_xg, y_pred_xg))
```

	precision	recall	f1-score	support
0	0.93	0.96	0.94	1049
1	0.70	0.57	0.63	184
accuracy			0.90	1233
macro avg	0.81	0.76	0.79	1233
weighted avg	0.89	0.90	0.90	1233

```
[50]: results = cross_val_score(modelxg, X, y, cv=kf)
      print("%0.2f accuracy with a standard deviation of %0.2f" % (results.
      ↪mean()*100, results.std()*100))
```

[17:50:34] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[17:50:36] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[17:50:38] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[17:50:40] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[17:50:41] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[17:50:43] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[17:50:45] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[17:50:47] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore

the old behavior.

[17:50:49] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[17:50:51] WARNING: ../src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

89.70 accuracy with a standard deviation of 0.87

After applying K-fold to the XGBoost model we can observe that the average score remains at 88-89%

[]: